
Semi-Supervised Factored Logistic Regression for High-Dimensional Neuroimaging Data

Anonymous Author(s)

Affiliation

Address

email

Abstract

Imaging neuroscience links human behavior to aspects of brain biology in ever-increasing datasets. Existing neuroimaging methods typically perform either discovery of unknown neural structure or testing of neural structure associated with mental tasks. However, testing hypotheses on the neural correlates of larger mental task sets necessitates adequate representations of the observations. We therefore propose to blend representation modelling and task classification into a unified statistical learning problem. A multinomial logistic regression is introduced that is constrained by factored coefficients and coupled with an autoencoder. We show that this approach yields more accurate and interpretable neural models of psychological tasks in a reference dataset, as well as better generalization to other datasets.

keywords: Brain Imaging, Cognitive Science, Semi-Supervised Learning, Systems Biology

1 Introduction

Methods for neuroimaging research can be grouped by discovering neurobiological structure or assessing the neural correlates associated with mental tasks. To *discover*, on the one hand, spatial distributions of neural activity structure across time, independent component analysis (ICA) is often used [?]. It decomposes the BOLD (blood-oxygen level-dependent) signals into the primary modes of variation. The ensuing spatial activity patterns are believed to represent brain networks of functionally interacting regions [?]. Similarly, sparse principal component analysis (SPCA) has been used to separate BOLD signals into parsimonious network components [?]. The extracted brain networks are probably manifestations of electrophysiological oscillation frequencies [?]. Their fundamental organizational role is further attested by continued covariation during sleep and anesthesia [?]. Network discovery by applying ICA or SPCA is typically performed on task-unrelated (i.e., *unlabeled*) “resting-state” data. These capture brain dynamics during ongoing random thought without controlled environmental stimulation. In fact, a large portion of the BOLD signal variation is known not to correlate with a particular behavior, stimulus, or experimental task [?].

To *test*, on the other hand, the neural correlates underlying mental tasks, the general linear model (GLM) is the dominant approach [?]. The contribution of individual brain voxels is estimated according to a design matrix of experimental tasks. Alternatively, psychophysiological interactions (PPI) elucidate the influence of one brain region on another conditioned by experimental tasks [?]. As a last example, an increasing number of neuroimaging studies model experimental tasks by training classification algorithms on brain signals [?]. All these methods are applied to task-associated (i.e., *labeled*) data that capture brain dynamics during stimulus-guided behavior. Two important conclusions can be drawn. First, the mentioned supervised neuroimaging analyses typically yield results in a voxel space. This ignores the fact that the BOLD signal exhibits spatially distributed

patterns of coherent neural activity. Second, existing supervised neuroimaging analyses cannot exploit the abundance of easily acquired resting-state data [?]. These may allow better discovery of the manifold of brain states due to the high task-rest similarities of neural activity patterns, as observed using ICA [?].

Both these neurobiological properties can be conjointly exploited in an approach that is mixed (i.e., using rest and task data), multitask and factored (i.e., performing network decomposition). The integration of brain-network discovery into supervised classification can yield a semi-supervised learning framework. The most relevant neurobiological structure should hence be identified for the prediction problem at hand. Autoencoders suggest themselves because they can emulate variants of most unsupervised learning algorithms, including PCA, SPCA, and ICA [?]. Autoencoders (AE) are often one-layered learning models that condense the input data to local and global representations via reconstruction under compression prior. They behave like a (truncated, non-centered) PCA in case of one linear hidden layer and a squared error loss [?]. Autoencoders behave like a SPCA if shrinkage terms are added to the model weights in the optimization objective. Moreover, they have the characteristics of an ICA in case of tied weights and adding a nonlinear convex function at the first layer [?]. These authors further demonstrated that ICA, sparse autoencoders, and sparse coding are mathematically equivalent under mild conditions. Thus, autoencoders may flexibly project the neuroimaging data onto the main directions of variation.

In the present investigation, a linear autoencoder will be fit to (unlabeled) rest data and integrated as a rank-reducing bottleneck into a multinomial logistic regression fit to (labeled) task data. We can then solve the unsupervised data representation and the supervised classification, previously studied in isolation, in an identical statistical learning problem. From the perspective of dictionary learning, the first layer represents projectors to the discovered set of basis functions which are linearly combined by the second layer to perform predictions[?]. Neurobiologically, this allows delineating a low-dimensional manifold of brain network patterns and then distinguishing mental tasks by their most discriminative linear combinations. Theoretically, a reduction in model variance should be achieved by resting-state autoencoders that privilege the most neurobiologically valid models in the hypothesis set. Practically, neuroimaging research is frequently suffering from data scarcity. This limits the set of representations that can be extracted from GLM analyses based on few participants. We therefore contribute a computational framework that 1) co-analyzes many problems simultaneously (thus finds shared representations by “multi-task learning”) and 2) exploits unlabeled data, as they span a space of meaningful configurations.

2 Methods

Data. As the currently biggest openly-accessible reference dataset, we chose the Human Connectome Project (HCP) resources [?]. Neuroimaging task data with labels of ongoing cognitive processes were drawn from 500 unrelated, healthy HCP participants. 18 HCP tasks were selected that are known to elicit reliable neural activity across participants. The task paradigms include 1) working memory/cognitive control processing, 2) incentive processing, 3) visual and somatosensory-motor processing, 4) language processing (semantic and phonological processing), 5) social cognition, 6) relational processing, and 7) emotional processing. All data were acquired on the same Siemens Skyra 3T scanner. Whole-brain EPI acquisitions were acquired with a 32 channel head




Figure 1: **Model architecture.** Linear autoencoders find an optimized compression of 79,941 brain voxels into n unknown activity patterns by improving reconstruction from them. The decomposition matrix equates with the bottleneck of a factored logistic regression. Supervised multi-class learning on task data (X_{task}) can thus be guided by unsupervised decomposition of rest data (X_{rest}).

coil (TR=720ms, TE=33.1ms, flip angle=52, BW=2290Hz/Px, in-plane FOV=280mm × 180mm, 72 slices, 2.0mm isotropic voxels). The “minimally preprocessed” pipeline includes gradient unwarping, motion correction, fieldmap-based EPI distortion correction, brain-boundary-based registration of EPI to structural T1-weighted scans, nonlinear (FNIRT) registration into MNI space, and grand-mean intensity normalization. Activity maps were spatially smoothed with a Gaussian kernel of 4mm (FWHM). A GLM was implemented by FILM from the FSL suite with model regressors from convolution with a canonical hemodynamic response function and from temporal derivatives. HCP tasks were conceived to modulate activity in a maximum of different brain regions and neural systems. Indeed, at least 70% of the participants showed consistent brain activity in contrasts from the task battery, which certifies excellent activity patterns covering extended parts of the brain [?]. In sum, the HCP task data incorporated 8650 first-level activity maps from 18 diverse paradigms administered to 498 participants (2 removed due to incomplete data). All maps were resampled to a common 60x72x60 space of 3mm isotropic voxels and gray-matter masked (at least 10% tissue probability). The supervised analyses were based on labeled HCP task maps with 79,941 voxels of interest representing z-values in gray matter.

These labeled data were complemented by unlabeled activity maps from HCP acquisitions of unconstrained resting-state activity [?]. These reflect brain activity in the absence of controlled thought. In line with the goal of the present study, acquisition of these data was specifically aimed at the study of task-rest correspondence. From each participant, we included two time-series for left and right phase encoding with 1,200 maps of multiband, gradient-echo planar imaging acquired during a period of 15min (TR=720 ms, TE=33.1 ms, flip angle=52, FOV=280mm × 180mm, and 2.0mm isotropic voxels). Besides run duration, the task acquisitions were identical to the resting-state fMRI acquisitions for maximal compatibility between task and rest data. We here drew on “minimally preprocessed” rest data from 200 randomly selected healthy, unrelated participants. PCA was applied to each set of 1,200 rest maps for denoising by keeping only the 20 main modes of variation. In sum, the HCP rest data concatenated 8000 unlabeled, noise-cleaned rest maps with 40 brain maps from each of 200 randomly selected participants.

We were further interested in the utility of the optimized low-rank projection in one task dataset for dimensionality reduction in another task dataset. To this end, the HCP-derived network decompositions were used as preliminary step in the classification problem of another large sample. The ARCH1 dataset [?] provides activity maps from diverse experimental tasks, including auditory and visual perception, motor action, reading, language comprehension and mental calculation. 81 right-handed healthy participants (3 not included in present analyses due to incomplete data) without psychiatric or neurological history participated in four fMRI sessions acquired under different experimental paradigms. The functional maps were warped into the MNI space and resampled to isotropic 3mm resolution. Whole-brain EPI data were acquired with the same Siemens Trio with a 32 channel head coil (TR=2400ms, TE=30ms, flip angle=60, in-plane FOV=192mm × 192mm, 40 slices, 3.0mm isotropic voxels). Standard preprocessing was performed with Nipype [?], including slice timing, motion correction, alignment, and spatial normalization. Activity maps were spatially smoothed by a Gaussian kernel of 5mm (FWHM). Analogous to HCP data, the second task dataset incorporated 1404 labeled, grey-matter masked, and z-scored activity maps from 18 diverse tasks acquired in 78 participants.

The labeled and unlabeled data were fed into a linear statistical model composed of an autoencoder and factored logistic regression.

Linear autoencoder. The affine autoencoder takes the input \mathbf{x} and projects it into coordinates of a latent representation \mathbf{z} and reconstructs it back to \mathbf{x}' by

$$\mathbf{z} = \mathbf{W}_0 \mathbf{x} + \mathbf{b}_0 \quad \mathbf{x}' = \mathbf{W}_1 \mathbf{z} + \mathbf{b}_1, \quad (1)$$

where $\mathbf{x} \in \mathbb{R}^d$ denotes the vector of $d = 79,941$ voxel values from each rest map, $\mathbf{z} \in \mathbb{R}^n$ is the n -dimensional hidden state (i.e., distributed neural activity patterns), and $\mathbf{x}' \in \mathbb{R}^d$ is the reconstruction vector of the original activity map from the hidden variables. Further, \mathbf{W}_0 denotes the weight matrix that transforms from input space into the hidden space (encoder), \mathbf{W}_1 is the weight matrix for back-projection from the hidden variables to the output space (decoder). \mathbf{b}_0 and \mathbf{b}_1 are bias vectors. The model parameters $\mathbf{W}_0, \mathbf{b}_0, \mathbf{b}_1$ are found by minimizing the expected squared reconstruction error

$$\mathbb{E} [\mathcal{L}_{\mathcal{AE}}(\mathbf{x})] = \mathbb{E} [\|\mathbf{x} - \mathbf{W}_1(\mathbf{W}_0 \mathbf{x} + \mathbf{b}_0) + \mathbf{b}_1\|^2]. \quad (2)$$

This reconstruction error criterion equates with maximizing a lower bound on the mutual information between input and the learned representation. Here we choose \mathbf{W}_0 and \mathbf{W}_1 to be tied, i.e. $\mathbf{W}_0 = \mathbf{W}_1^T$. In particular, this means that the weights learned are forced to take a two-fold function: That of signal *analysis* and that of signal *synthesis*. The first layer *analyzes* the data in order to obtain the cleanest latent representation, while the second layer represents building blocks from which to *synthesize* the data using the latent activations. Tying these processes together makes the analysis layer interpretable and pulls all non-zero singular values towards 1. Nonlinearities were not applied to the activations in the first layer.

Factored logistic regression. Lossy compression by a low-dimensional bottleneck is also imposed by the first layer of the factored multinomial logistic regression. It gives the probability of an input \mathbf{x} to belong to a class $i \in \{1, \dots, l\}$

$$P(Y = i | \mathbf{x}; \mathbf{V}_0, \mathbf{V}_1, \mathbf{c}_0, \mathbf{c}_1) = \text{softmax}_i(f_{\mathcal{LR}}(\mathbf{x})), \quad (3)$$

where $f_{\mathcal{LR}}(\mathbf{x}) = \mathbf{V}_1(\mathbf{V}_0\mathbf{x} + \mathbf{c}_0) + \mathbf{c}_1$ computes multinomial logits and $\text{softmax}_i(x) = \exp(x_i) / \sum_j \exp(x_j)$. The matrix $\mathbf{V}_0 \in \mathbb{R}^{d \times n}$ transforms the input $\mathbf{x} \in \mathbb{R}^d$ into n latent components and the matrix $\mathbf{V}_1 \in \mathbb{R}^{n \times l}$ projects the latent components onto hyperplanes that reflect l label probabilities. \mathbf{c}_0 and \mathbf{c}_1 are corresponding bias vectors. The loss function is given by

$$\mathbb{E}[\mathcal{L}_{\mathcal{LR}}(\mathbf{x}, \mathbf{y})] \approx \frac{1}{N_{X_{task}}} \sum_{k=0}^{N_{X_{task}}} \log(P(Y = y^{(k)} | \mathbf{x}^{(k)}; \mathbf{V}_0, \mathbf{V}_1, \mathbf{c}_0, \mathbf{c}_1)). \quad (4)$$

Layer combination. Importantly, the optimization problem of the linear autoencoder and the factored logistic regression are linked on two levels. First, their transformation matrices mapping from input to the latent space are tied

$$\mathbf{V}_0 = \mathbf{W}_0. \quad (5)$$

We thus search for a compression of the 79,941 voxel values into n latent components that represent a latent code optimized for both rest and task activity data. Second, the objectives of the autoencoder and the factored logistic regression are interpolated in the common loss function

$$\mathcal{L}(\theta, \lambda) = \lambda \mathcal{L}_{\mathcal{LR}} + (1 - \lambda) \frac{1}{N_{X_{rest}}} \mathcal{L}_{\mathcal{AE}} + \Omega(\theta). \quad (6)$$

In so doing, we search for the combined model parameters $\theta = \{\mathbf{V}_0, \mathbf{V}_1, \mathbf{c}_0, \mathbf{c}_1, \mathbf{b}_0, \mathbf{b}_1\}$ with respect to the (unsupervised) reconstruction error and the (supervised) task classification. $\mathcal{L}_{\mathcal{AE}}$ is divided by $N_{X_{rest}}$ to equilibrate both loss terms to the same order of magnitude. $\Omega(\theta)$ represents a regularization, for which we choose the ElasticNet-type, i.e., a combination of ℓ_1 and ℓ_2 penalties.

Optimization. The common objective was optimized by gradient descent in the SS-FLogReg parameters. The required gradients were obtained by using the chain rule to backpropagate error derivatives. We chose the *rmsprop* solver [?], a refinement of stochastic gradient descent. *Rmsprop* dictates an adaptive learning rate for each model parameter by scaled gradients from a running average. The batch size was set to 100 (much expected redundancy in X_{rest} and X_{task}), matrix parameters were initialized by Gaussian random values multiplied by 0.004 (i.e., gain), and bias parameters were initialized to zero. The normalization factor and the update rule for θ are given by

$$\begin{aligned} \mathbf{v}^{(t+1)} &= \rho \mathbf{v}^{(t)} + (1 - \rho) \left(\nabla_{\theta} f(x^{(t)}, y^{(t)}, \theta^{(t)}) \right)^2 \\ \theta^{(t+1)} &= \theta^{(t)} + \alpha \frac{\nabla_{\theta} f(x^{(t)}, y^{(t)}, \theta^{(t)})}{\sqrt{\mathbf{v}^{(t+1)} + \epsilon}}, \end{aligned} \quad (7)$$

where f is the loss function computed on a minibatch sample at timestep t , α is the learning rate (0.00001), ϵ a global damping factor (10^{-6}) and ρ the decay rate (0.9 to deemphasize the magnitude of the gradient). Note that we have also experimented with other solvers (stochastic gradient descent, adadelta, and adagrad) but found that *rmsprop* converged faster and with similar or higher generalization performance.

Hints. In fact, the constraint by a rest-data autoencoder qualifies as a *hint* rather than regularization in a strict sense [?]. Its purpose is to introduce prior knowledge on *known* properties of the *unknown* target function f . Rather than only relying on input-output pairs in the learning process, we thus narrow our hypothesis set to the biologically most plausible solutions. That is, we reduce the search space in a way that is compatible with the expected representation of BOLD activity signals.

Implementation. The analyses were performed in Python. We used *nilearn* to handle the large quantities of neuroimaging data [?] and *Theano* for automatic, numerically stable differentiation of symbolic computation graphs [?, ?]. All Python scripts that generated the results are accessible online for reproducibility and reuse <http://github.com/anonymous/anonymous>.

3 Experimental Results

Serial versus parallel structure discovery and classification. We first tested whether there is a substantial advantage in combining unsupervised decomposition and supervised classification learning. We benchmarked our approach against performing data reduction on the (unlabeled) first half of the HCP task data by PCA, SPCA, ICA, and AE ($n = 5, 20, 50, 100$ components) and learning classification models in the (labeled) second half by ordinary logistic regression. PCA reduced the dimensionality of the task data by finding orthogonal network components using a change of basis (whitening). SPCA separated the BOLD signals into network components with few regions by a regression-type optimization problem constrained by ℓ_1 -penalty (no orthogonality assumptions, 1000 maximum iterations, per-iteration tolerance of $1 * 10^{-8}$, $\alpha = 1$). ICA performed iterative blind source separation by a parallel FASTICA implementation (200 maximum iterations, per-iteration tolerance of 0.0001, initialized by random mixing matrix, whitening). AE found a code of latent representations by optimizing projection into a bottleneck (500 iterations, same implementation as below for rest data). The second half of the task data was projected onto the latent components discovered in its first half. Only the ensuing component loadings were submitted to ordinary logistic regression (one hidden layer, $\ell_1 = 0.1$, $\ell_2 = 0.1$, 500 maximum iterations). These serial two-step approaches were compared against parallel decomposition and classification by FLogReg (two hidden layers, $\lambda = 1.00$, $\ell_1 = 0.1$, $\ell_2 = 0.1$, 500 iterations). Importantly, all trained classification models were tested on a large, unseen test set (20% of data) in the present analyses. Across choices for n , FLogReg achieved more than 95% out-of-sample accuracy, whereas supervised learning based on PCA, SPCA, ICA, and AE loadings ranged from 32% to 87% (Table 1). This experiment establishes the advantage of directly searching for classification-relevant structure in the fMRI data, rather than solving the supervised and unsupervised problems independently. This effect was particularly pronounced when assuming few hidden dimensions.

| n | PCA + LogReg | SPCA + LogReg | ICA + LogReg | AE + LogReg | FLogReg |
|-----|--------------|---------------|--------------|-------------|--------------|
| 5 | 45,1 % | 32,2 % | 37,5 % | 44,2 % | 95,7% |
| 20 | 78,1 % | 78,2 % | 81,0 % | 63,2 % | 97,3% |
| 50 | 81,7 % | 84,0 % | 84,2 % | 77,0 % | 97,6% |
| 100 | 81,3 % | 82,2 % | 87,3 % | 76,6 % | 97,4% |

Table 1: **Serial versus parallel dimensionality reduction and classification.** Chance is at 5,6%.

Model performance. SS-FLogReg was subsequently trained (500 epochs) across parameter choices for the hidden components ($n = 5, 20, 100$) and the balance between autoencoder and logistic regression ($\lambda = 0, 0.25, 0.5, 0.75, 1$). Assuming 5 latent directions of variation should yield models with higher bias and smaller variance than SS-FLogReg with 100 latent directions. Given the 18-class problem of HCP, setting λ to 0 consistently yields generalization performance at chance-level (5,6%) because only the unsupervised layer of the estimator is optimized. At each epoch (i.e., iteration over the data), the out-of-sample performance of the trained classifier was assessed on 20% of unseen HCP data. Additionally, the “out-of-study performance” of the learned decomposition was assessed by using it as dimensionality reduction of an independent labeled dataset (i.e., ARCH1) and conducting ordinary logistic regression on the ensuing component loadings.

We made several noteworthy observations (Table 2). First, the purely supervised estimator ($\lambda = 1$) achieved in no instance the best accuracy, precision, recall or f1 scores on HCP data. Classification by factored logistic regression is therefore facilitated by imposing structure from the unlabeled rest

| | $n = 5$ | | | | | $n = 20$ | | | | | $n = 100$ | | | | |
|------------------------|---------------|------------------|-----------------|------------------|---------------|---------------|------------------|-----------------|------------------|---------------|---------------|------------------|-----------------|------------------|---------------|
| | $\lambda = 0$ | $\lambda = 0.25$ | $\lambda = 0.5$ | $\lambda = 0.75$ | $\lambda = 1$ | $\lambda = 0$ | $\lambda = 0.25$ | $\lambda = 0.5$ | $\lambda = 0.75$ | $\lambda = 1$ | $\lambda = 0$ | $\lambda = 0.25$ | $\lambda = 0.5$ | $\lambda = 0.75$ | $\lambda = 1$ |
| Out-of-sample accuracy | 6,0% | 88,9% | 95,1% | 96,5% | 95,7% | 5,5% | 97,4% | 97,8% | 97,3% | 97,3% | 6,1% | 97,2% | 97,0% | 97,8% | 97,4% |
| Precision (mean) | 5,9% | 87,0% | 94,9% | 96,3% | 95,4% | 5,1% | 97,4% | 97,1% | 97,0% | 97,0% | 5,9% | 96,9% | 96,5% | 97,5% | 96,9% |
| Recall (mean) | 5,6% | 88,3% | 95,2% | 96,6% | 95,7% | 4,6% | 97,5% | 97,5% | 97,4% | 97,4% | 7,2% | 97,2% | 97,2% | 97,9% | 97,4% |
| F1 score (mean) | 4,1% | 86,6% | 94,9% | 96,4% | 95,4% | 3,8% | 97,4% | 97,2% | 97,1% | 97,1% | 5,3% | 97,0% | 96,7% | 97,7% | 97,2% |
| Out-of-study accuracy | 39,4% | 60,8% | 54,3% | 60,7% | 62,9% | 77,0% | 79,7% | 81,9% | 79,7% | 79,4% | 79,2% | 82,2% | 81,7% | 81,3% | 75,8% |

Table 2: Performance of SS-FLogReg across model parameter choices. Chance is at 5,6%.

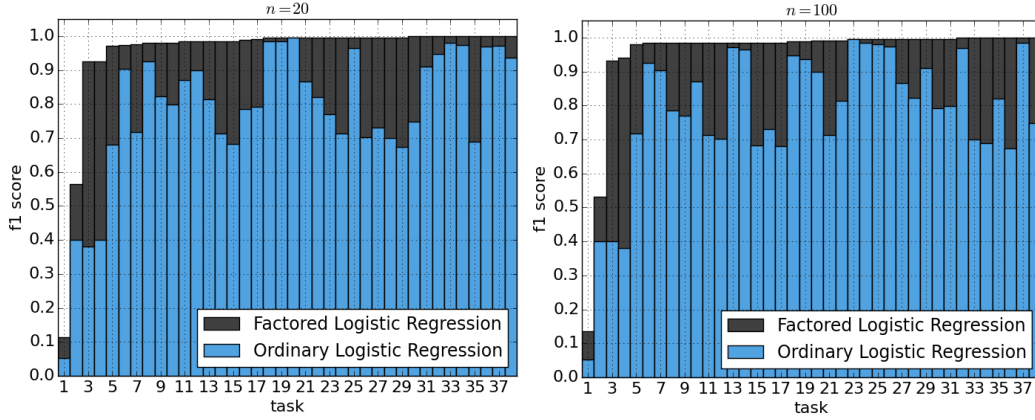


Figure 2: Effect of bottleneck in a 38-task classification problem. Depicts the f1 score for each of 38 class predictions (chance is 2,6%) sorted from lowest to highest. Multinomial logistic regression operating in voxel space (blue bars) was compared to factored logistic regression operating in 20 (left plot) and 100 (right plot) latent modes (grey bars). Autoencoder or rest data were not used for these analyses ($\lambda = 1$). Ordinary logistic regression yielded 77,7% accuracy out of sample, while factored logistic regression scored at 94,4% ($n = 20$) and 94,2% ($n = 100$). Hence, compressing the voxel data into a component space for classification achieves higher task separability.

data. Second, the higher the number of latent components n , the higher the out-of-study performance with small values of λ . This suggests that the presence of more rest-data-inspired hidden components results in more effective feature representations in unrelated task data. Third, for $n = 20$ and 100 (but not 5) the purely rest-data-trained decomposition matrix ($\lambda = 0$) resulted in noninferior out-of-study performance of 77,0% and 79,2%, respectively (Table 2). This confirms that guiding model learning by task-unrelated structure extracts features of general relevance beyond the supervised problem at hand.

Individual effects of dimensionality reduction and rest data. We first quantified the impact of introducing a bottleneck layer that disregards the autoencoder. To this end, ordinary logistic regression was juxtaposed with FLogReg at $\lambda = 1$ (no autoencoder). For this experiment, we increased the difficulty of the classification problem by including data from all 38 HCP tasks. Indeed, increased class separability in component space, as compared to voxel space, entails differences in generalization performance of $\approx 17\%$ (Figure 2). Notably, the cognitive tasks on reward and punishment processing are among the least predicted with ordinary but well predicted with factored logistic regression (tasks 3 and 4 in Figure 2). These experimental conditions have indeed been reported to exhibit highly similar neural activity patterns in GLM analyses of that dataset [?]. Consequently, also local activity differences (in the striatum and visual cortex in this case) can be successfully captured by brain-network modelling.

We then quantified the impact of structure from rest data keeping all model parameters constant ($n = 20$, $\lambda = 0.5$, $\ell_1 = 0.1$, $\ell_2 = 0.1$). At the beginning of every epoch, 2000 task and 2000 rest maps were drawn with replacement from same amounts of task maps but varying amounts of rest maps. In frequently encountered data-scarce settings (≈ 100 samples), but not in data-rich settings (≈ 1000 samples), the repertoire of rest structure modulated model performance (Figure 3).

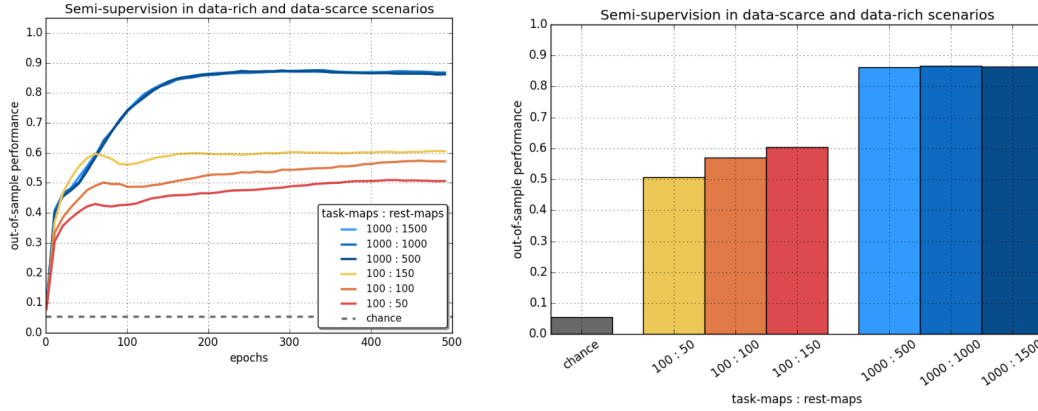


Figure 3: **Effect of richness in rest data** Gradient descent was performed on 2000 task and 2000 rest maps. These were drawn with replacement from identical versus varying quantities of task versus rest maps, at the beginning of each epoch. Chance is at 5,6%.

Feature identification. We finally examined whether the models were fit for purpose (Figure 4). To this end, we computed Pearson’s correlation between the classifier weights and the averaged neural activity map for each of the 18 tasks. Ordinary logistic regression thus yielded a mean correlation of $\rho = 0.28$. For SS-FLogReg ($\lambda = 0.25, 0.5, 0.75, 1$), a per-class-weight map was computed by matrix multiplication of the two inner layers. Feature identification thus ranged between $\rho = 0.35$ and $\rho = 0.55$ for $n = 5$, between $\rho = 0.59$ and $\rho = 0.69$ for $n = 20$, and between $\rho = 0.58$ and $\rho = 0.69$ for $n = 100$. Consequently, SS-FLogReg puts higher absolute weights on relevant structure. This reflects a higher signal-to-noise ratio, in part explained by the more BOLD-typical local contiguity. Conversely, SS-FLogReg puts lower probability mass on irrelevant structure. Despite lower interpretability, the salt-and-pepper-like weight maps obtained from ordinary logistic regression are sufficient for good classification performance. Hence, SS-FLogReg yielded class weights that were much more similar to features of the respective training samples for any choice of λ . SS-FLogReg therefore captures genuine properties of neural activity patterns, rather than participant- or study-specific artefacts.

Miscellaneous observations. For the sake of completeness, we informally report modifications of the model that did not improve the generalization performance. Introducing stochasticity into model learning by a) input corruption of \mathbf{X}_{task} or b) drop-out at the bottleneck \mathbf{W}_0 using binomial distributions ($p = 0.1, 0.3, 0.5$) deteriorated model performance in all instances. Adding c) rectified linear units (ReLU) to \mathbf{W}_0 or other commonly used nonlinearities (d) sigmoid, e) soft-plus, f) hyperbolic tangent) all led to smaller classification accuracies, probably due to sample size limits. Further, g) “pretraining” of the bottleneck \mathbf{W}_0 (i.e., non-random initialization) by either corresponding SPCA or ICA loadings did not exhibit improved accuracies, neither did h) autoencoder pretraining. Moreover, introducing an additional i) overcomplete layer (100 units) after the bottleneck was not advantageous. Finally, imposing either j) only ℓ_1 -penalty or k) only ℓ_2 -penalty was disadvantageous in all tested cases, which favored ElasticNet regularization chosen in the above analyses.

4 Discussion and Conclusion

Using the flexibility of factored models, we learn the low-dimensional representation from high-dimensional voxel brain space that is most important for prediction of cognitive task sets. From a machine-learning perspective, factorization of the logistic regression weights can be viewed as transforming a “multi-class classification problem” into a “multi-task learning problem”. The higher generalization accuracy and feature recovery, comparing to ordinary logistic regression, hold potential for adoption in various neuroimaging analyses. Besides increased performance, these models are more interpretable by automatically learning a mapping to and from a brain-network space. This domain-specific learning algorithm encourages departure from the artificial and statistically less at-

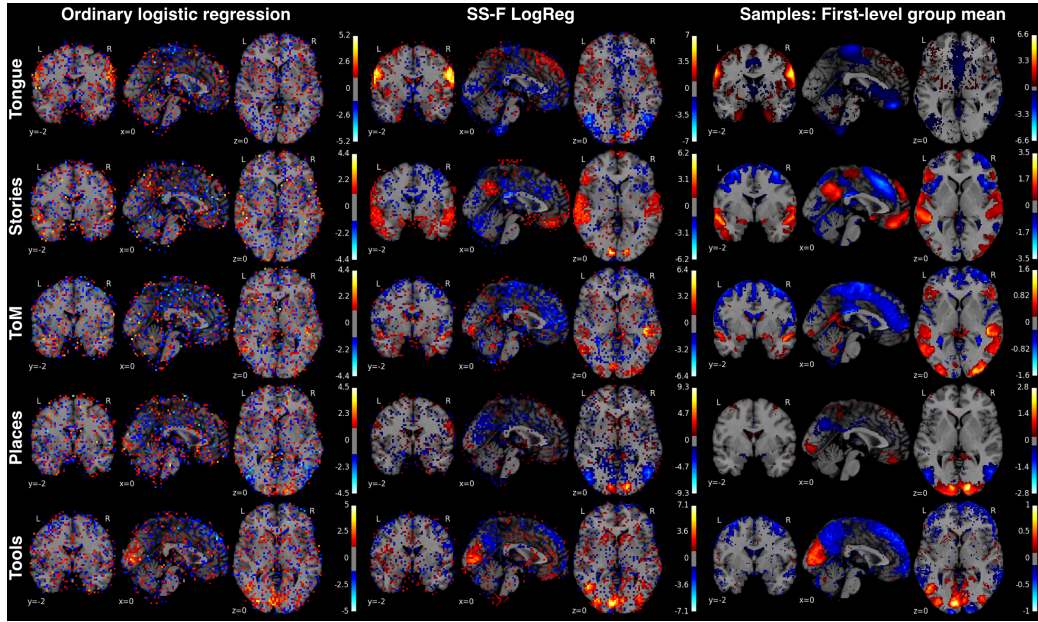


Figure 4: **Classification weight maps** The voxel predictors corresponding to 5 exemplary (of 18 total) psychological tasks (*rows*) from the HCP dataset [?]. *Left column*: multinomial logistic regression (same implementation but without bottleneck or autoencoder), *middle column*: semi-supervised factored logistic regression ($n = 20$ latent components, $\lambda = 0.5$, $\ell_1 = 0.1$, $\ell_2 = 0.1$), *right column*: voxel-wise average of whole-brain activity maps from each task. Semi-supervised factored logistic regression a) puts higher absolute weights on relevant structure, b) lower ones on irrelevant structure, and c) yields BOLD-typical local contiguity (without enforcing a spatial prior). All values are z-scored and thresholded at the 75th percentile.

432 tractive voxel space. Neurobiologically, brain activity underlying defined mental operations can be
433 explained by linear combinations of the main activity patterns. That is, fMRI data probably con-
434 centrate near a low-dimensional manifold of characteristic brain networks combinations. Extracting
435 fundamental building blocks of brain organization might facilitate the quest for the cognitive prim-
436 itives of human thought. We hope that these first steps stimulate development towards powerful
437 semi-supervised representation extraction in systems neuroscience.

438 In the future, automatic reduction of brain maps to their neurobiological essence may leverage data-
439 intense neuroimaging studies. Initiatives for data collection are rapidly increasing in neuroscience
440 [?]. These promise structured integration of neuroscientific knowledge accumulating in databases.
441 Tractability by condensed feature representations can avoid the ill-posed problem of learning the full
442 distribution of activity patterns. This is not only relevant to the multi-class challenges spanning the
443 human cognitive space [?] but also the multi-modal combination with high-resolution 3D models of
444 brain anatomy [?] and high-throughput genome analyses [?]. The biggest socioeconomic potential
445 may lie in across-hospital clinical studies that predict disease trajectories and drug responses in
446 psychiatric/neurological populations [?, ?].

447 **Acknowledgment** Data were provided by the Human Connectome Project.
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485